

# From Classical Domain Decomposition Solvers to Learning

---

**Victorita Dolean**

with: A. Heinlein, S. Mishra, B. Moseley

LJLL French-Spanish summer school, 7-11 July 2025



A. Heinlein, B. Moseley, S. Mishra

## Artificial Neural Networks for Solving Ordinary and Partial Differential Equations

Isaac Elias Lagaris, Aristidis Likas, *Member, IEEE*, and Dimitrios I. Fotiadis

Published in **IEEE TRANSACTIONS ON NEURAL NETWORKS, VOL. 9, NO. 5, 1998.**

**Idea:** Solve a general differential equation subject to boundary conditions

$$G(x, \Psi(x), \nabla \Psi(x), \nabla^2 \Psi(x)) = 0 \quad \text{in } \Omega$$

by solving an **optimization problem**

$$\min_p \sum_{x_i} G(x_i, \Psi_t(x_i, p), \nabla \Psi_t(x_i, p), \nabla^2 \Psi_t(x_i, p))^2$$

where  $\Psi_t(x, p)$  is a **trial function**,  $x_i$  **sampling points inside the domain**  $\Omega$  and  $p$  are **adjustable parameters**.

# Main Features of the Method

## Construction of the trial functions

The trial functions **explicitly satisfy the boundary conditions**:

$$\Psi_t(\mathbf{x}, \mathbf{p}) = A(\mathbf{x}) + F(\mathbf{x}, N(\mathbf{x}, \mathbf{p})),$$

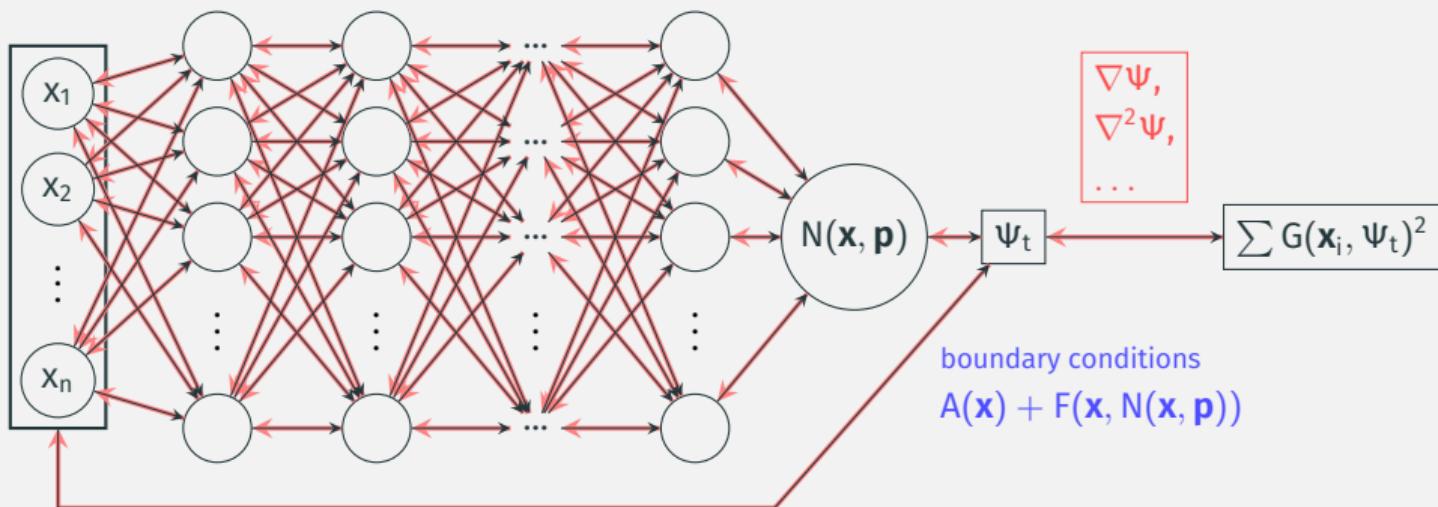
where

- $N(\mathbf{x}, \mathbf{p})$  is a **trainable feedforward neural network** with parameters  $\mathbf{p}$  and input  $\mathbf{x} \in \mathbb{R}^n$  and
- the functions  $A$  and  $F$  are **fixed functions** chosen such that
  - $A$  **satisfies the boundary conditions**, and
  - $F$  **does not contribute to the boundary conditions**.

From the conclusion of the paper:

“The **success of the method** can be attributed to **two factors**. The first is the employment of neural networks that are **excellent function approximators** and the second is the form of the trial **solution that satisfies by construction the BC's** and therefore the constrained optimization problem becomes a substantially simpler unconstrained one.”

## Sketch of the Approach by Lagaris et al. (1998)



# Physics-Informed Neural Networks (PINNs)

In **Raissi, Perdikaris, Karniadakis (2019)**, the authors have revisited and modified the approach by **Lagaris et al. (1998)**, denoting their method as **physics-informed neural networks (PINNs)**.

Consider the partial differential equation

$$\mathcal{N}[u](\mathbf{x}, t) = f(\mathbf{x}, t), \quad (\mathbf{x}, t) \in [0, T] \times \Omega \subset \mathbb{R}^d.$$

The main novelty of PINNs is that a **hybrid loss function** is used for training the feedforward neural network:

$$\mathcal{L} = \omega_{\text{data}} \mathcal{L}_{\text{data}} + \omega_{\text{PDE}} \mathcal{L}_{\text{PDE}},$$

where  $\omega_{\text{data}}$  and  $\omega_{\text{PDE}}$  are **weights** and

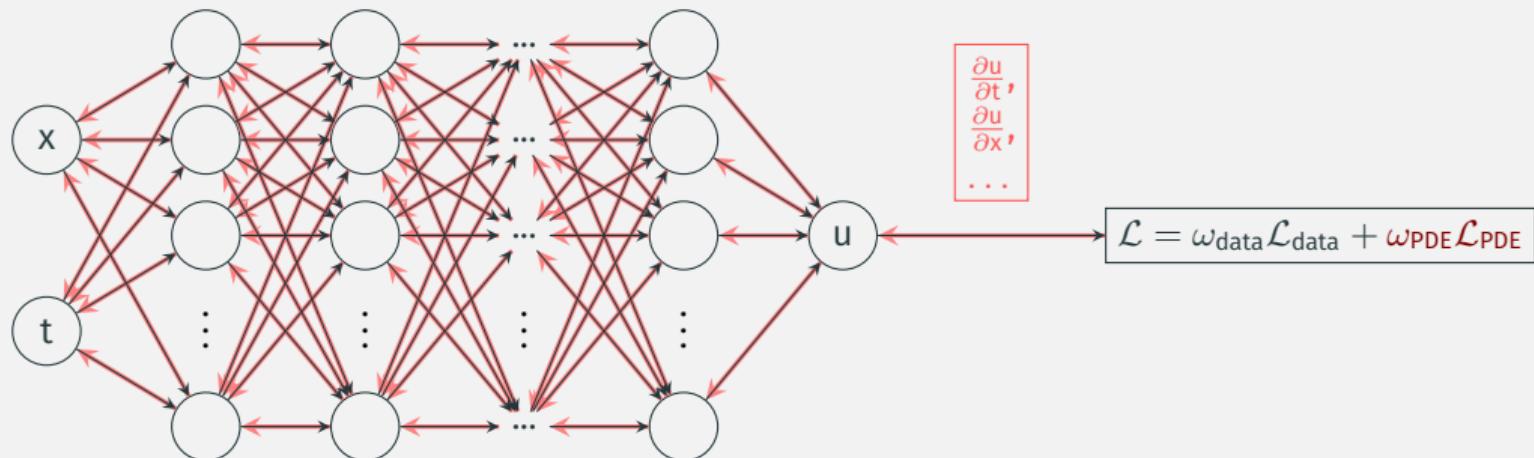
$$\mathcal{L}_{\text{data}} = \frac{1}{N_{\text{data}}} \sum_{i=1}^{N_{\text{data}}} (u(\mathbf{x}_i, t_i) - u_i)^2,$$

$$\mathcal{L}_{\text{PDE}} = \frac{1}{N_{\text{PDE}}} \sum_{i=1}^{N_{\text{data}}} (\mathcal{N}[u](\mathbf{x}_i, t_i) - f(\mathbf{x}_i, t_i))^2.$$



- **Known solution values** can be included in  $\mathcal{L}_{\text{data}}$ .
- **IC and BC** are also included in  $\mathcal{L}_{\text{data}}$

## Sketch of the PINN approach by Raissi et al.



## Available theoretical results: PINNs

Mishra, S. and Molinaro, R. Estimates on the generalisation error of PINNs, 2022

### Estimate of the generalisation error

The generalisation error (or total error) verifies

$$\mathcal{E}_G \leq C_{\text{pde}} \mathcal{E}_T + C_{\text{pde}} C_{\text{quad}}^{1/p} N^{-\alpha/p}$$

where

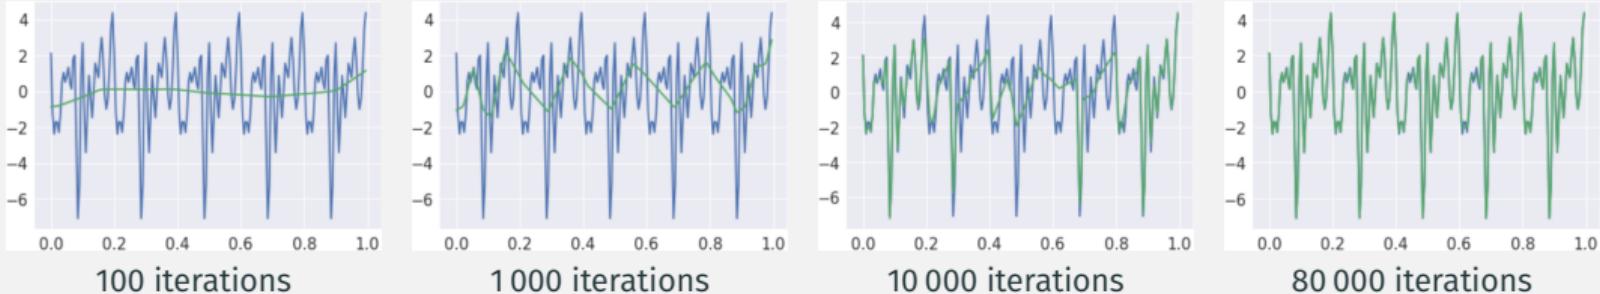
- $\mathcal{E}_G = \mathcal{E}_G(\theta^*; \mathcal{S}) := \|\mathbf{u} - \mathbf{u}^*\|_X$
- $\mathcal{E}_T$  is the training error
- $C_{\text{pde}}$  and  $C_{\text{quad}}$ : constants depending on the PDE and the quadrature
- $N$ : number of the training points

The devil is in the details:

**“As long as the PINN is trained well it generalises well”**

# Scaling Issues in Neural Network Training

- **Spectral bias:** Neural networks prioritise learning lower frequency functions first irrespective of their amplitude.



Rahaman, N., et al, On the spectral bias of neural networks, **ICML (2019)**

Hong, Siegel, Tan, Xu, On the Activation Function Dependence of the Spectral Bias of Neural Networks, **arXiv (2022)**

- Solving solutions on large domains and/or with multiscale features potentially requires **very large neural networks**.
- Training may **not sufficiently reduce the loss** or take **large numbers of iterations**.
- Significant **increase on the computational work**

# When PINNs fail to train?

Perdikaris, P. et al, When and why PINNs fail to train: A neural tangent kernel perspective, 2022.

## Neural tangent kernel (NTK) theory in a nutshell

- Write the gradient descent method at the continuous level ( $\mathcal{L}$  is the loss function)

$$\frac{d\theta}{dt} = -\nabla \mathcal{L}(\theta), \quad \mathcal{L}(\theta) := \omega_b \sum \mathcal{R}(\mathbf{x}_b, \theta(t))^2 + \omega_r \sum \mathcal{R}(\mathbf{x}_r, \theta(t))^2$$

- Residual vectors in the collocation points obey an ODE

$$\frac{d}{dt} \begin{bmatrix} \mathcal{R}(\mathbf{x}_b, \theta(t)) \\ \mathcal{R}(\mathbf{x}_r, \theta(t)) \end{bmatrix} = -\mathbf{K}(t) \begin{bmatrix} \mathcal{R}(\mathbf{x}_b, \theta(t)) \\ \mathcal{R}(\mathbf{x}_r, \theta(t)) \end{bmatrix}$$

- The Neural Tangent Kernel  $\mathbf{K}(t) \rightarrow \mathbf{K}^*$  for infinitely wide and shallow networks
- Spectral properties of  $\mathbf{K}^*$  explain the speed of training.

**"To provide further insight, we analyze the training dynamics of fully-connected PINNs through the lens of their NTK and show that not only they suffer from spectral bias, but they also exhibit a discrepancy in the convergence rate among the different loss components contributing to the total training error"**

# A Motivation for Domain Decomposition Approaches

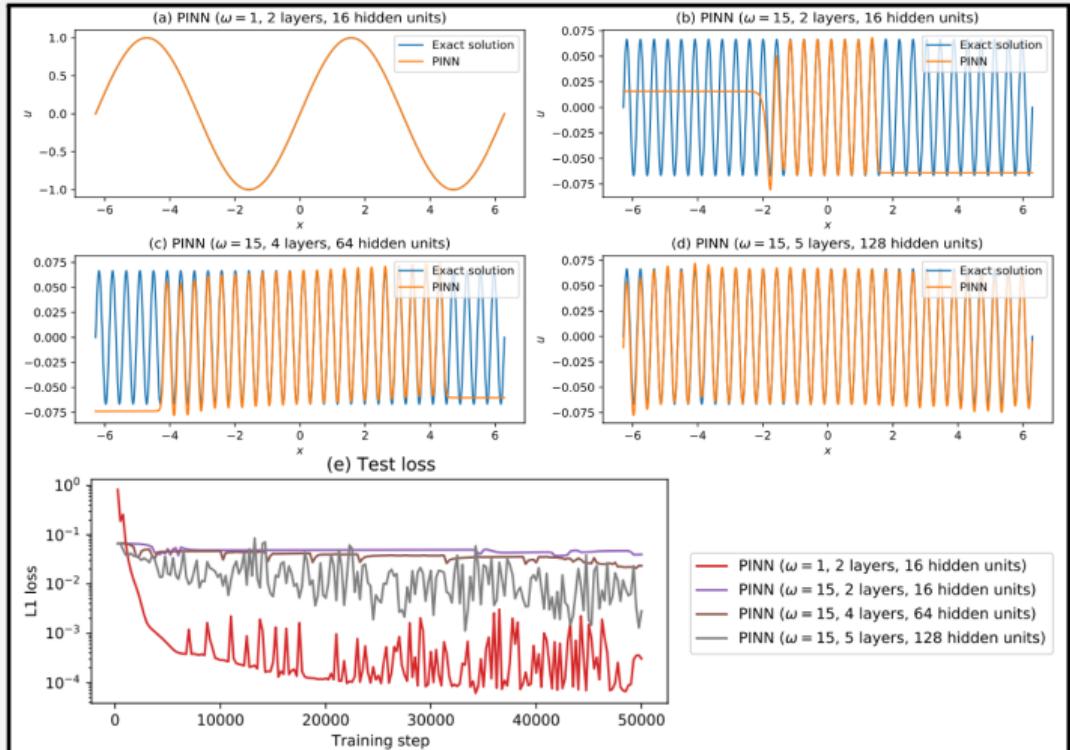
Solve

$$\begin{aligned} u' &= \cos(\omega x), \\ u(0) &= 0, \end{aligned}$$

for different values of  $\omega$ .

## Scaling issues

- Size of the computational domain
- Size of frequencies



(a) 321 free parameters

(d) 66 433 free parameters

# Finite Basis Physics-Informed Neural Networks (FBPINNs)

In the **finite basis physics informed neural network (FBPINNs) method** introduced in **Moseley, Markham, and Nissen-Meyer (arXiv 2021)**, we solve the boundary value problem

$$\begin{aligned}\mathcal{N}[u](\mathbf{x}) &= f(\mathbf{x}), & \mathbf{x} \in \Omega \subset \mathbb{R}^d, \\ \mathcal{B}_k[u](\mathbf{x}) &= g_k(\mathbf{x}), & \mathbf{x} \in \Gamma_k \subset \partial\Omega.\end{aligned}$$

using neural networks, we employ the **PINN** approach and **enforce the boundary conditions using a constraining operator**, similar to **Lagaris et al. (1998)**.

## Weak enforcement of boundary conditions

Loss function

$$\mathcal{L}(\theta) = w_I \mathcal{L}_{\text{PDE}} + w_B \mathcal{L}_{\text{BC}},$$

where

$$\begin{aligned}\mathcal{L}_{\text{PDE}}(\theta) &= \frac{1}{N_I} \sum_{i=1}^{N_I} (\mathcal{N}[u](\mathbf{x}_i, \theta) - f(\mathbf{x}_i))^2, \\ \mathcal{L}_{\text{BC}}(\theta) &= \frac{1}{N_B} \sum_{i=1}^{N_B} (\mathcal{B}_k[u](\mathbf{x}_i, \theta) - g_k(\mathbf{x}_i))^2.\end{aligned}$$

## Hard enforcement of boundary conditions

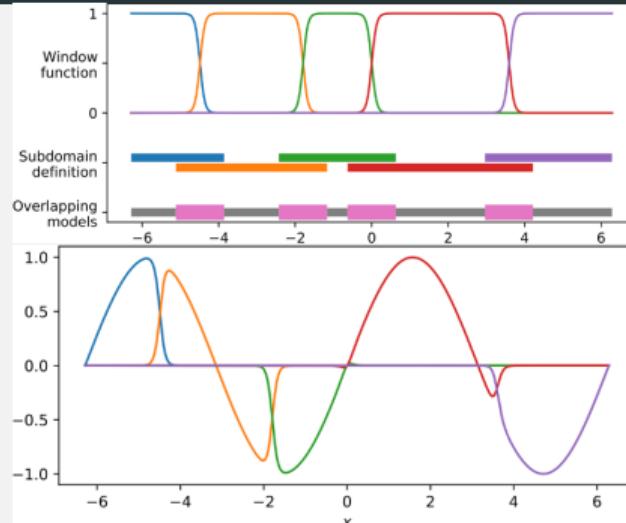
Loss function

$$\mathcal{L}(\theta) = \frac{1}{2} \sum_{i=1}^{N_I} (\mathcal{N}[\mathcal{C}u](\mathbf{x}_i, \theta) - f(\mathbf{x}_i))^2,$$

with constraining operator  $\mathcal{C}$ , which explicitly enforces the boundary conditions.

→ Often **improves training performance**

# FBPINNs – Overlapping Domain Decomposition



- Domain decomposition:  $\Omega = \bigcup_{j=1}^J \Omega_j$
- Collocation points (global):  $\{x_i\}_{i=1}^N$
- Overlapping/interior parts  $\Omega_j^\circ$  and  $\Omega_j^{\text{int}}$
- Local solutions  $u_j$ , window functions  $w_j$
- Global solution  $\mathcal{C}u = \mathcal{C} \sum_{j, x_i \in \Omega_j} \omega_j u_j$

Global loss function

$$\mathcal{L}(\theta_1, \dots, \theta_J) = \underbrace{\frac{1}{N} \sum_{j=1}^J \sum_{x_i \in X_j^\circ} \left( \mathcal{N}[\mathcal{C} \sum_{l, x_i \in X_l} \omega_l u_l](x_i, \theta_l) - f(x_i) \right)^2}_{=: \mathcal{L}^\circ(\theta_1, \dots, \theta_J)} + \underbrace{\frac{1}{N} \sum_{j=1}^J \sum_{x_i \in X_j^{\text{int}}} \left( \mathcal{N}[\mathcal{C} \sum_{l, x_i \in X_l} \omega_l u_l](x_i, \theta_l) - f(x_i) \right)^2}_{=: \mathcal{L}^{\text{int}}(\theta_1, \dots, \theta_J)}.$$

Since  $X_i^{\text{int}} \cap X_j^{\text{int}} = \emptyset$  for  $i \neq j$ ,

$$\mathcal{L}^{\text{int}}(\theta_1, \dots, \theta_J) = \frac{1}{N} \sum_{j=1}^J \sum_{x_i \in X_j^{\text{int}}} (\mathcal{N}[\mathcal{C} \omega_j u_j](x_i, \theta_j) - f(x_i))^2$$

- The subdomains can be split **active** (trained in parallel) and **inactive** (fixed)
- This corresponds to classical **parallel (all active)** or **multiplicative (one active at a time)** Schwarz methods

# FBPINN With Flexible Scheduling

FBPINN training step

**if**  $j \in \mathcal{A}$  ( $\Omega_j$  is an active domain) **then**

Perform  $p$  iterations of gradient descent on  $\theta_j$  ( $\theta_l^k, l \neq j$  are kept fixed):

$$\theta_j^{k+l} = \theta_j^{k+l-1} - \lambda \nabla_{\theta_j} \mathcal{L}(\theta_1^k, \dots, \theta_j^{k+l-1}, \dots, \theta_p^k), l = 1, \dots, p.$$

Update the solutions in the overlap (communication with the neighbours):

$$\forall \mathbf{x} \in \Omega_j^\circ, u(\mathbf{x}, \theta_j^{k+p}) \leftarrow \sum_{l, \mathbf{x} \in \Omega_l} \omega_l u_l(\mathbf{x}, \theta_l^{k+p}).$$

Update the gradients in the overlap.

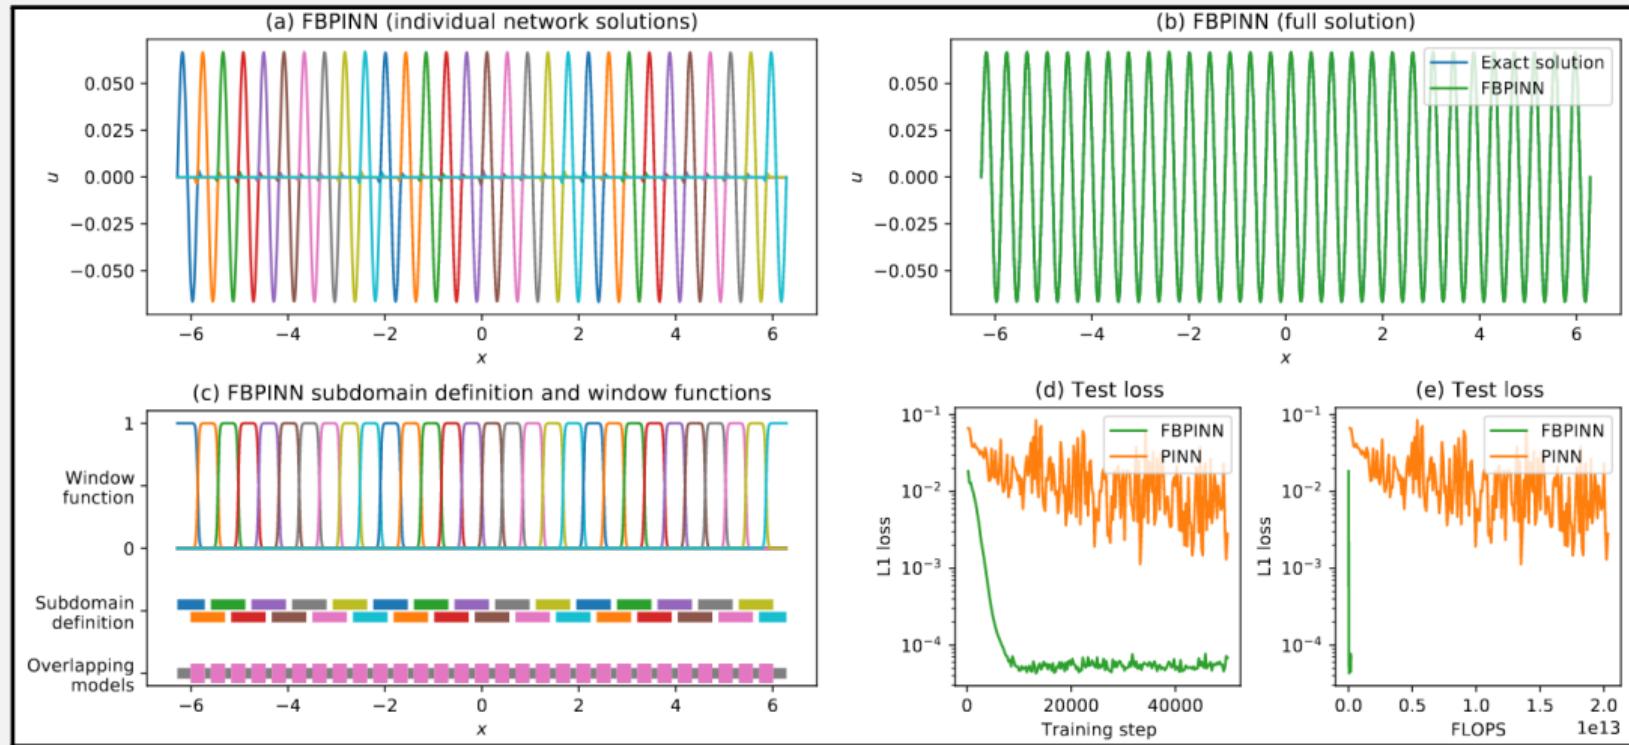
Apply the constraining operator  $u(\mathbf{x}, \theta_j^k) \leftarrow \mathcal{C}(u(\mathbf{x}, \theta_j^{k+p})).$

**end if**

## Summary

- Communication every  $p$  iterations (better for overall efficiency)
- Multiplication with a window function : a way to restrict to the local domain.
- The set of active domains can be changed after the local training is completed.

# Numerical Results – Comparison of PINNs and FBPINNs



# Two-Level FBPINN Algorithm

## Coarse correction and spectral bias

### Questions:

- Scalability requires **global transport of information**. This can be done via **coarse global problem**.
- What does this mean in the **context of network training**?

### Idea:

- Learn low frequencies using a small **global network**, train high frequencies using local networks.

## Two-level FBPINN network architecture:

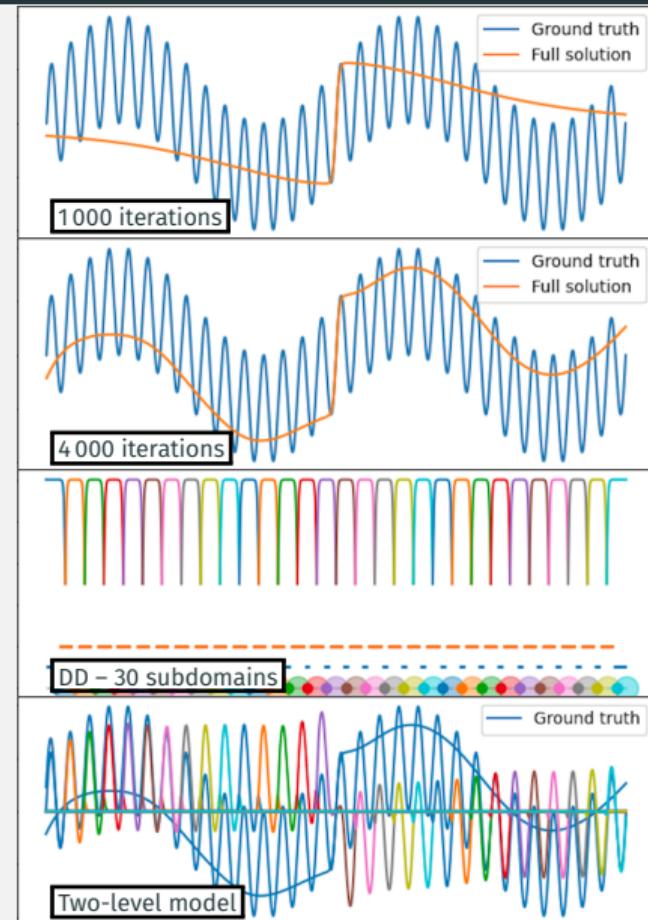
$$u(\theta_0, \theta_1, \dots, \theta_l) = \mathcal{C}(u_0(\theta_l) + \sum_{\Omega_j} \omega_l u_l(\theta_l))$$

Consider a **simple model problem** with **two frequencies**

$$\begin{cases} \frac{du}{dx} = \omega_1 \cos(\omega_1 x) + \omega_2 \cos(\omega_2 x) \\ u(0) = 0. \end{cases}$$

with  $\omega_1 = 1, \omega_2 = 15$ .

Cf. **V. Dolean, A. Heinlein, S. Mishra, B. Moseley**, Finite basis physics-informed neural networks as a Schwarz domain decomposition method, [arXiv:2211.05560, 2022](https://arxiv.org/abs/2211.05560)



# Laplace Problem – One- Vs Two-Level FBPINN

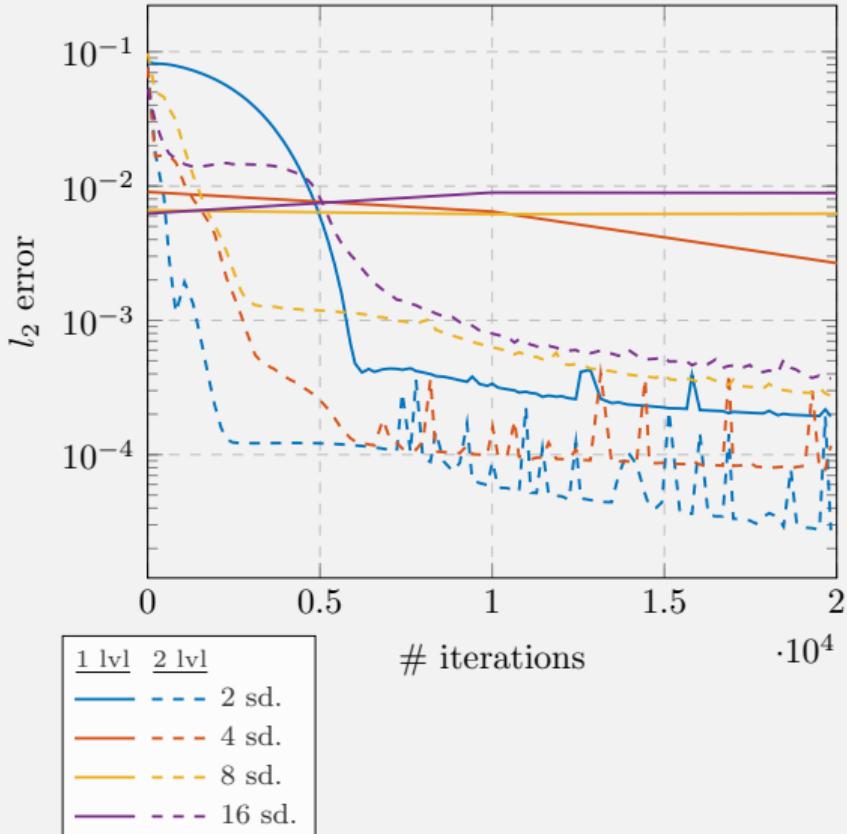
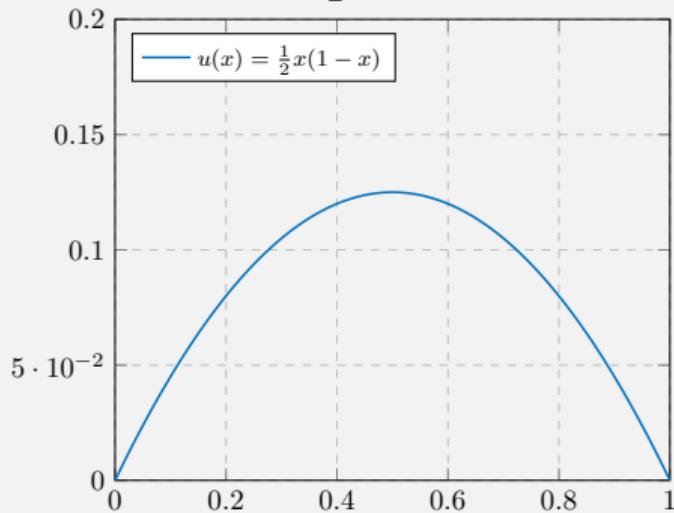
Let us now consider the **simple boundary value problem**

$$-u'' = 1 \quad \text{in } [0, 1],$$

$$u(0) = u(1) = 0,$$

which has the **solution**

$$u(x) = \frac{1}{2}x(1-x).$$



# Multi-Frequency Laplace Problem – One- Vs Two-Level FBPINN

## Multi-frequency boundary value problem

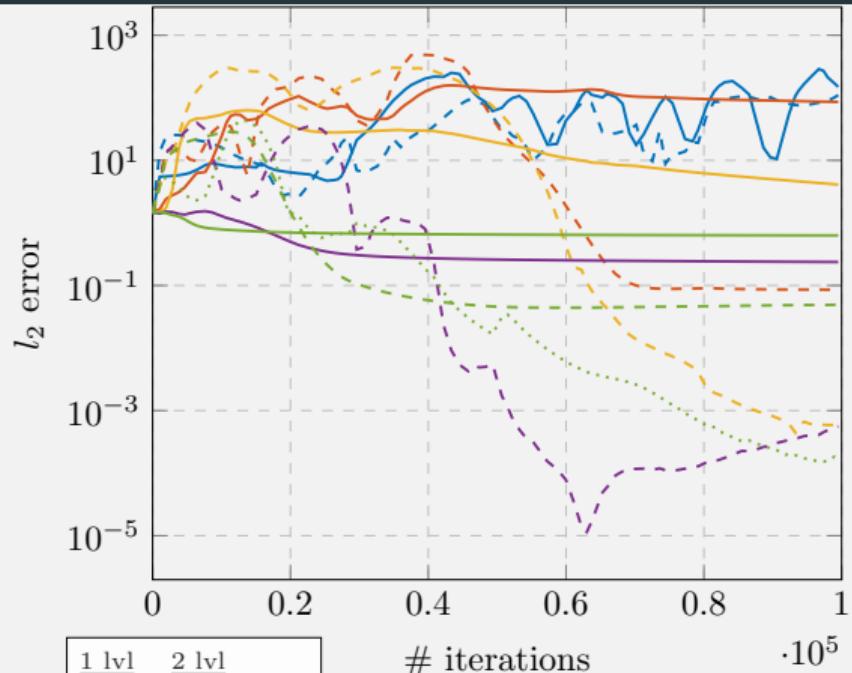
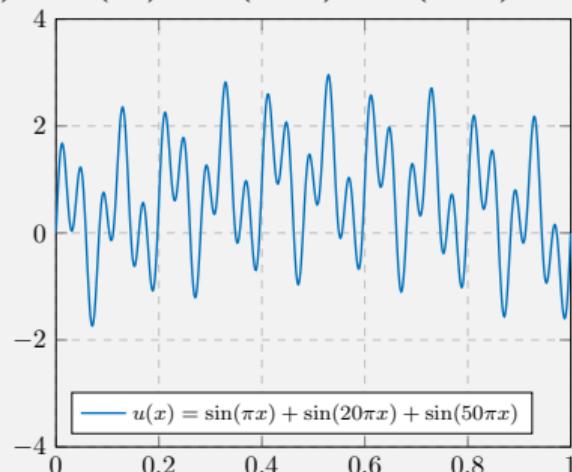
$$-u'' = \pi^2 \sin(\pi x) + (20\pi)^2 \sin(20\pi x)$$

$$+ (50\pi)^2 \sin(50\pi x) \text{ in } [0, 1],$$

$$u(0) = u(1) = 0,$$

## Solution

$$u(x) = \sin(\pi x) + \sin(20\pi x) + \sin(50\pi x).$$



1 lvl	2 lvl
— 2 sd	- - - 2 sd
— 4 sd	- - - 4 sd
— 8 sd	- - - 8 sd
— 16 sd	- - - 16 sd
— 32 sd	- - - 32 sd
.... 2 coarse sd	- - - 2 coarse sd

## Conclusions

- Accurate approximation requires large numbers of subdomains
- Coarse level enables scalability

# Multilevel FBPINN Algorithm

## L-level network architecture

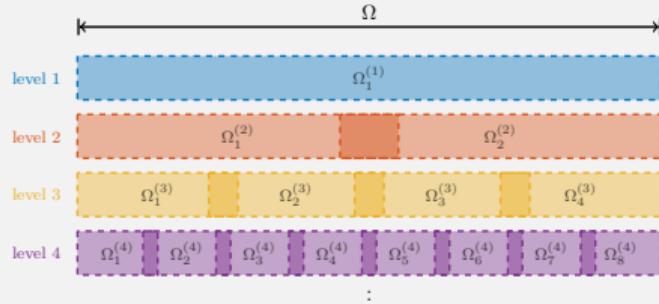
Let  $L$  overlapping domain decompositions

$\Omega = \bigcup_{j=1}^{J^{(l)}} \Omega_j^{(l)}$  and **window functions**  $\omega_j^{(l)}$  with

$$\sum_{j=1}^{J^{(l)}} \omega_j^{(l)} \equiv 1, \text{ supp}(\omega_j^{(l)}) \subset \Omega_j^{(l)}.$$

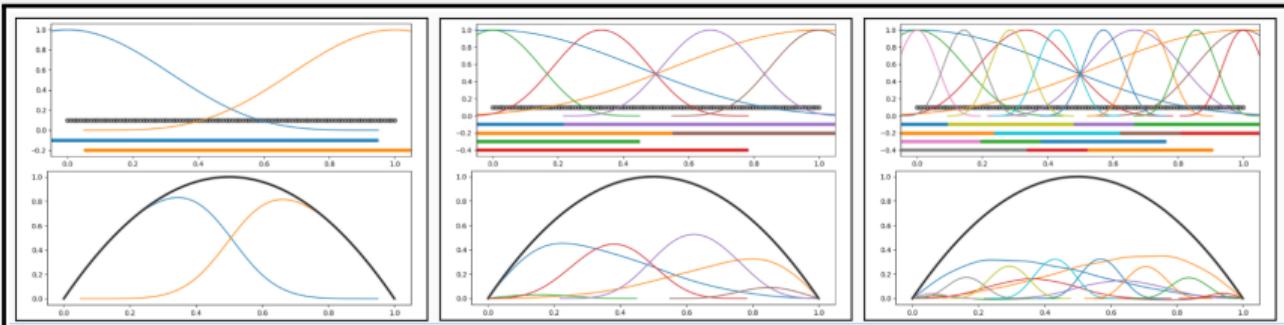
Then the multilevel architecture is

$$u(\mathbf{x}, \theta) = \mathcal{C} \sum_{l=1}^L \sum_{j=1}^{J^{(l)}} \omega_j^{(l)} u_j^{(l)}(\mathbf{x}, \theta_j^{(l)})$$



## Loss function

$$\mathcal{L}(\theta) = \frac{1}{N} \sum_{i=1}^N (\mathcal{N}[\mathcal{C} \sum_{l=1}^L \sum_{\mathbf{x}_i \in \Omega_j^{(l)}} \omega_j^{(l)} u_j^{(l)}](\mathbf{x}_i, \theta_j^{(l)}) - f(\mathbf{x}_i))^2.$$



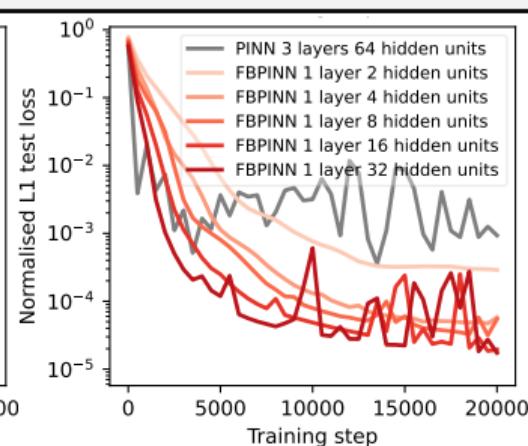
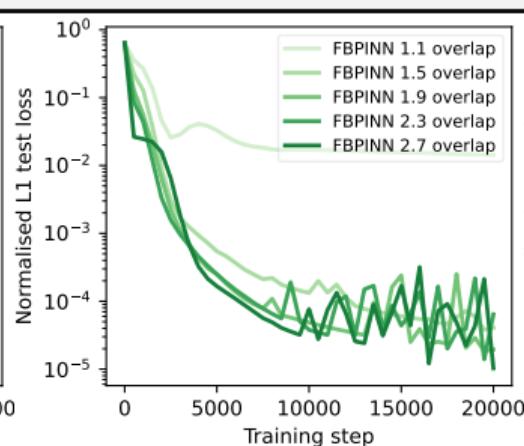
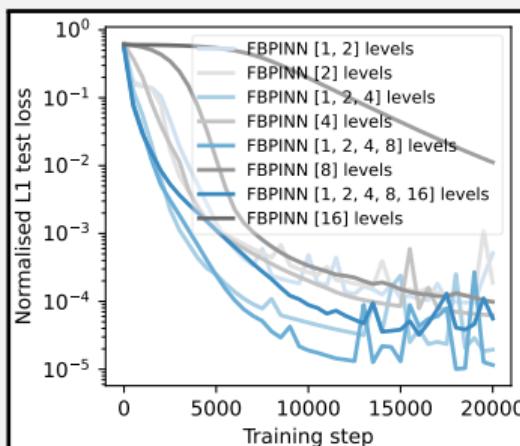
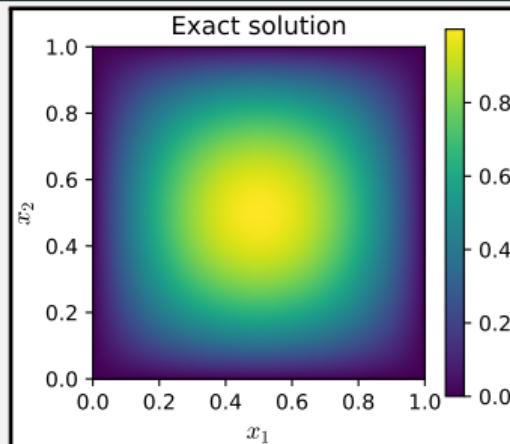
## Multilevel FBPINNs – 2D Laplace

Let us consider the **simple two-dimensional boundary value problem**

$$\begin{aligned} -\Delta u &= 32(x(1-x) + y(1-y)) && \text{in } \Omega = [0, 1]^2, \\ u &= 0 && \text{on } \partial\Omega, \end{aligned}$$

which has the **solution**

$$u(x, y) = 16(x(1-x)y(1-y)).$$



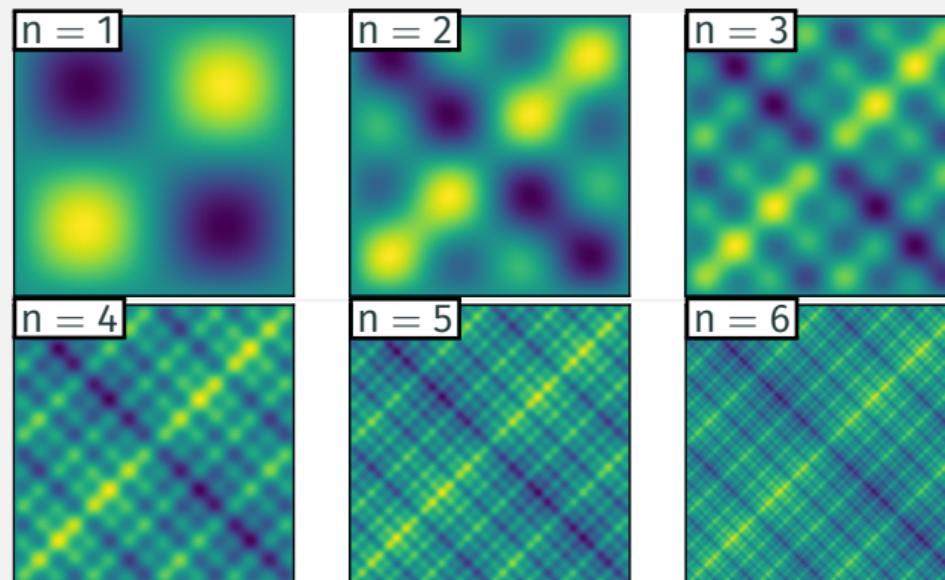
## Multi-Frequency Problem

Let us now consider the **two-dimensional multi-frequency Laplace boundary value problem**

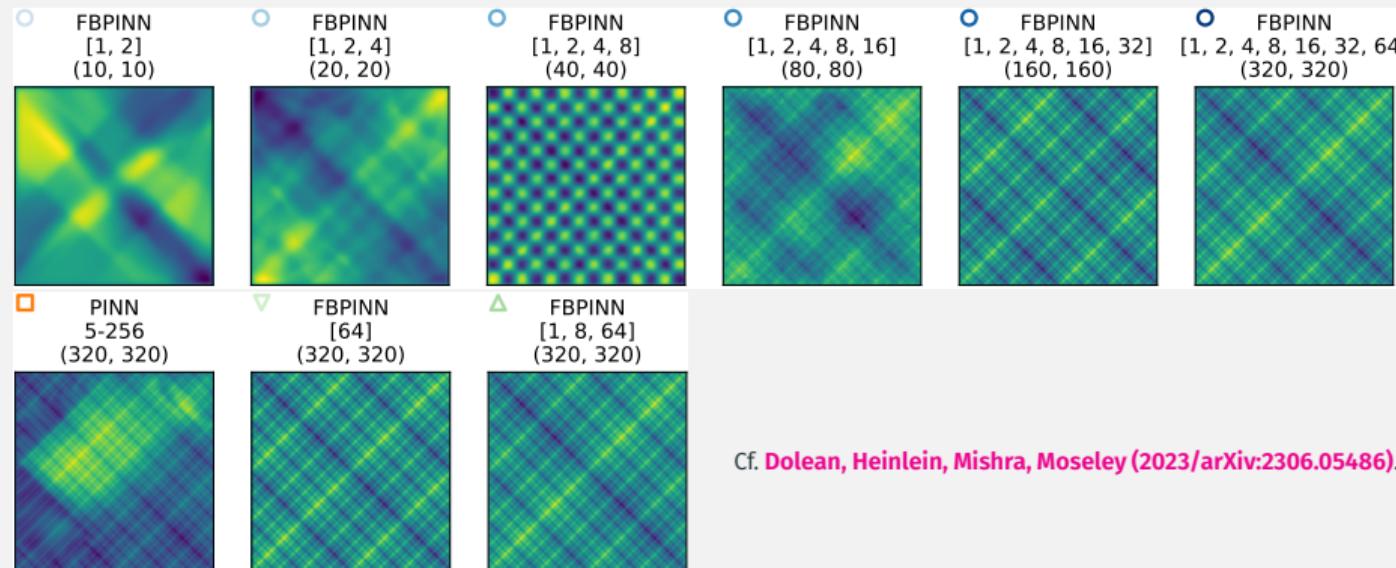
$$\begin{aligned} -\Delta u &= 2 \sum_{i=1}^n (\omega_i \pi)^2 \sin(\omega_i \pi x) \sin(\omega_i \pi y) && \text{in } \Omega = [0, 1]^2, \\ u &= 0 && \text{on } \partial\Omega, \end{aligned}$$

with  $\omega_i = 2^i$ .

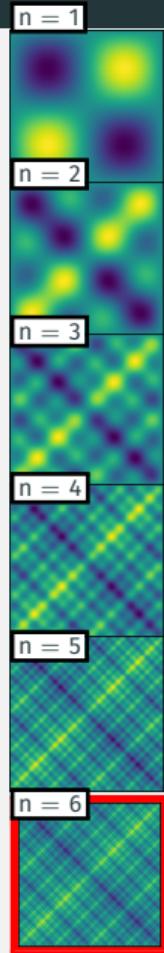
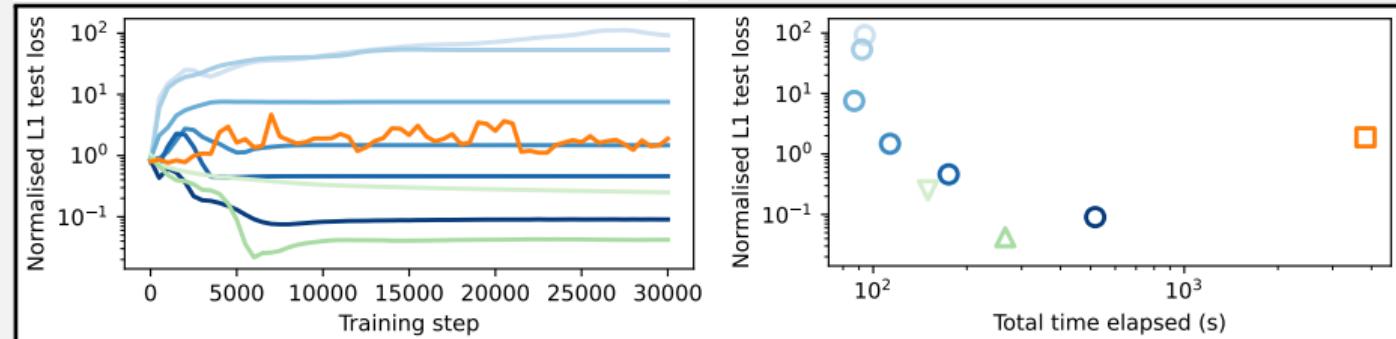
For increasing values of  $n$ , we obtain the **analytical solutions**:



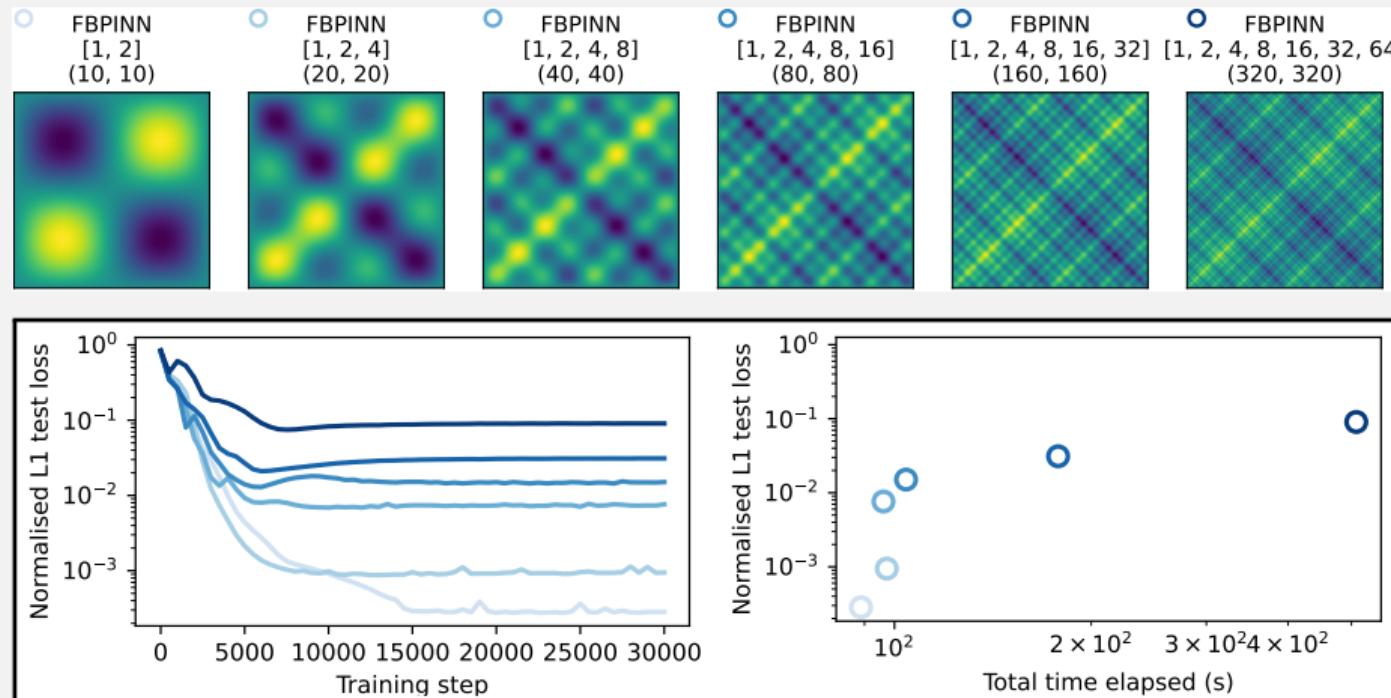
# Multi-Level FBPINNs for a Multi-Frequency Problem – Strong Scaling



Cf. [Dolean, Heinlein, Mishra, Moseley \(2023/arXiv:2306.05486\)](https://arxiv.org/abs/2306.05486).



# Multi-Level FBPINNs for a Multi-Frequency Problem – Weak Scaling



Cf. Dolean, Heinlein, Mishra, Moseley (2023/arXiv:2306.05486).

## Helmholtz Problem

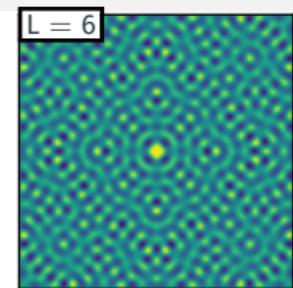
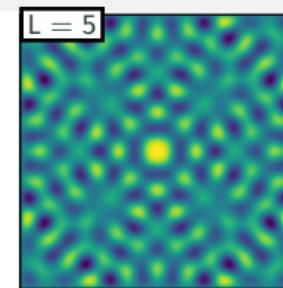
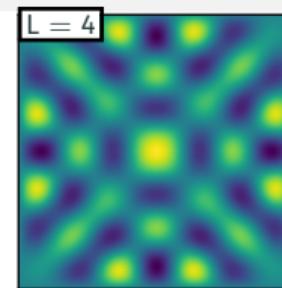
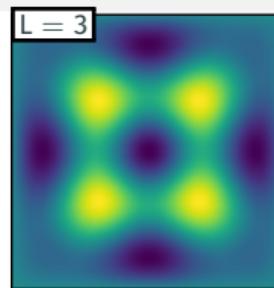
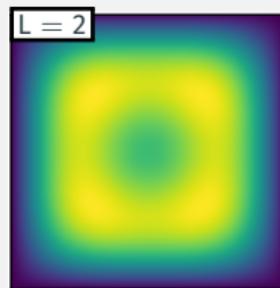
Finally, let us consider the **two-dimensional Helmholtz boundary value problem**

$$\Delta u - k^2 u = f \quad \text{in } \Omega = [0, 1]^2,$$

$$u = 0 \quad \text{on } \partial\Omega,$$

$$f(\mathbf{x}) = e^{-\frac{1}{2}(\|\mathbf{x} - 0.5\|/\sigma)^2}.$$

With  $k = 2^L \pi / 1.6$  and  $\sigma = 0.8 / 2^L$ , we obtain the **solutions**:



# Multilevel FBPINNs – 2D Helmholtz Problem

Let us consider the **two-dimensional Helmholtz boundary value problem**

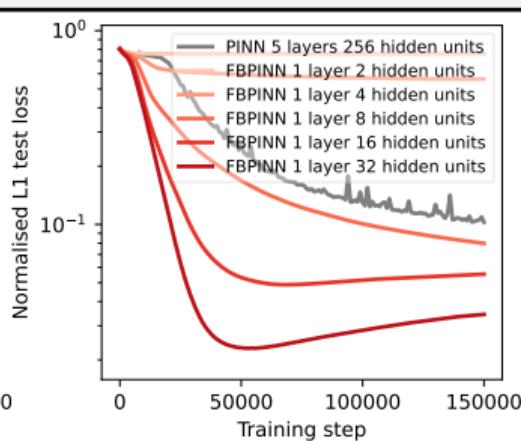
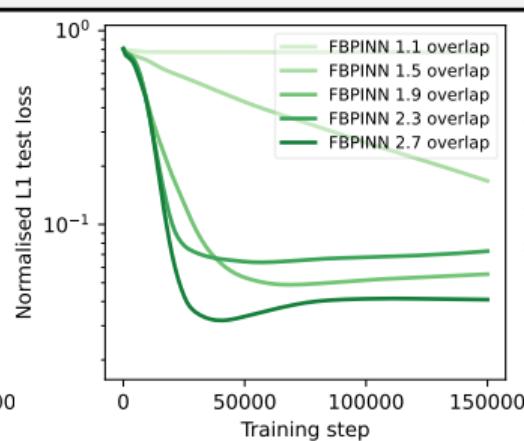
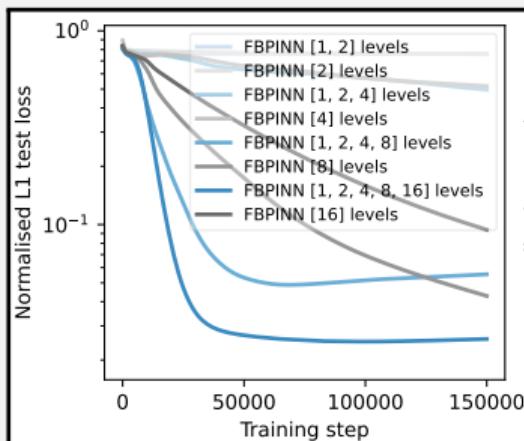
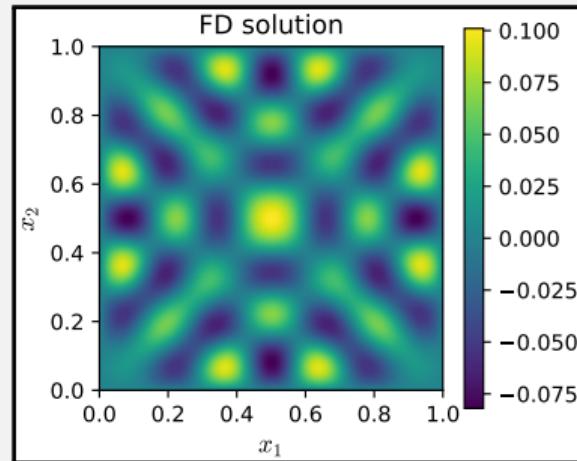
$$\Delta u - k^2 u = f \quad \text{in } \Omega = [0, 1]^2,$$

$$u = 0 \quad \text{on } \partial\Omega,$$

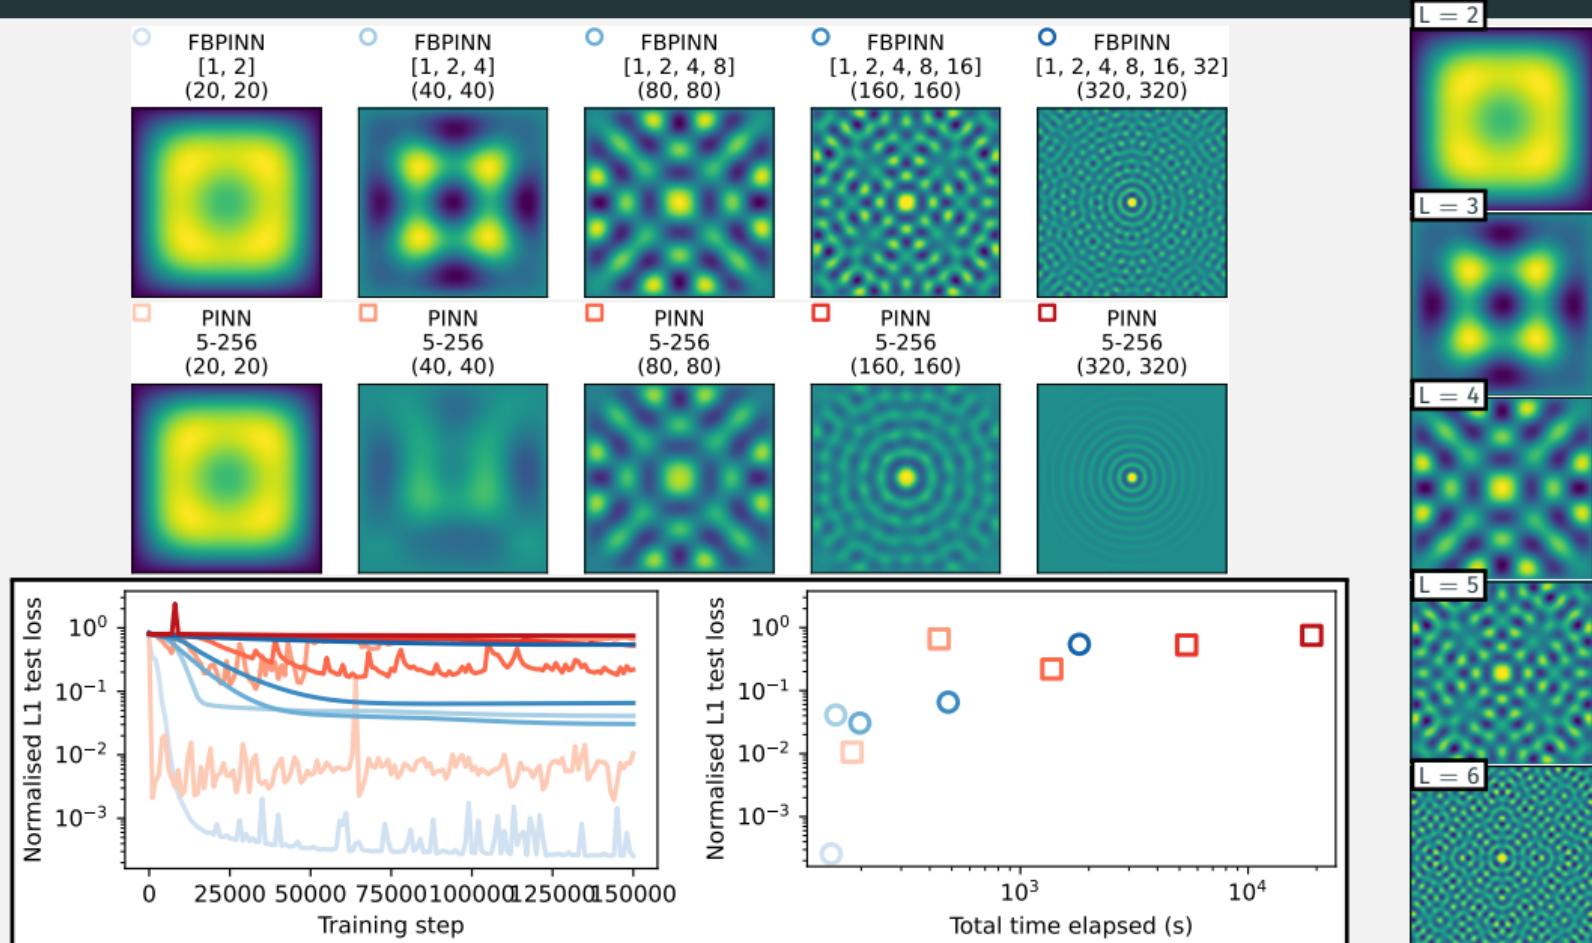
$$f(x) = e^{-\frac{1}{2}(\|x-0.5\|/\sigma)^2}.$$

with  $k = 2^4\pi/1.6$  and  $\sigma = 0.8/2^4$ .

We compute a **reference solution** using **finite differences** with a 5-point stencil on a  $320 \times 320$  grid.



# Multi-Level FBPINNs for the Helmholtz Problem – Weak Scaling



## Conclusions

- Interpretation of FBPINNs from a DD perspective brings the fields of ML and classical DD closer together GitHub: <https://github.com/benmoseley/FBPINNs>
- Multilevel versions have a potential for high wave numbers in an inverse problem setting.
- The multilevel training idea not limited to PINNs or a given architecture.
- **BUT** training times are still more important than traditional methods when solving PDEs ↵ multi-GPU training, parallel FBPINN training algorithm, more inputs to the subdomains networks ... more theory?

Thanks for your attention!